

Semantic Spotter Project

Index

1. Problem Statement
 - 1.1 Background
 - 1.2 Core Problem
 - 1.3 System Requirements
2. System Architecture
 - 2.1 System Architecture Overview
 - 2.2 Architectural Layers
3. Implementation
4. A Flowchart Explaining the System Design and the Various Layers
 - 3.1 System Flowchart Overview
 - 3.2 Layer-wise Flow Explanation
5. Detailed Description of the Problem Statement and Why LangChain / LlamaIndex is an Ideal Framework
 - 4.1 Detailed Description of the Problem Statement
 - 4.2 Why LangChain / LlamaIndex is an Ideal Framework
 - 4.2.1 Modular Design
 - 4.2.2 Native RAG Support
 - 4.2.3 Seamless Vector Store Integration
 - 4.2.4 LLM Flexibility
6. Innovation and Creativity in System Design
 - 5.1 Fully Local RAG Pipeline
 - 5.2 Optimized Chunking Strategy
 - 5.3 Prompt Engineering for Grounding
 - 5.4 Modular Architecture
7. Optimum System Architecture, Workflow and Implementation
 - 6.1 System Architecture Layers
 - 6.1.1 Data Ingestion Layer
 - 6.1.2 Preprocessing Layer
 - 6.1.3 Embedding Layer
 - 6.1.4 Vector Storage Layer
 - 6.1.5 Retrieval Layer
 - 6.1.6 LLM Layer
 - 6.1.7 Orchestration Layer
 - 6.2 Workflow
 - 6.3 Implementation Efficiency
8. Appropriate Use of the Components of LangChain / LlamaIndex in the System Design
 - 7.1 Document Loader

- 7.2 Text Splitter**
- 7.3 Embedding Model**
- 7.4 Vector Store**
- 7.5 Retriever**
- 7.6 Prompt Template**
- 7.7 Chain Builder**

- 9. Detailing Project Goals, Data Sources, Design Choices and Challenges Faced**

- 8.1 Project Goals**
- 8.2 Data Sources**
- 8.3 Design Choices**
- 8.4 Challenges Faced**

- 10. Conclusion**

problem statement

Organizations deal with large volumes of unstructured documents such as insurance policies, legal agreements, contracts, and operational manuals. Extracting precise information from these documents manually is time-consuming, inefficient, and error-prone.

Traditional keyword-based search systems fail because:

- They do not understand semantic meaning.
- They cannot interpret natural language questions.
- They often retrieve irrelevant results.

The core problem addressed in this project is:

How can we design an intelligent system that understands user queries in natural language and retrieves the most contextually relevant information from large PDF documents, while minimizing hallucinations?

The system must:

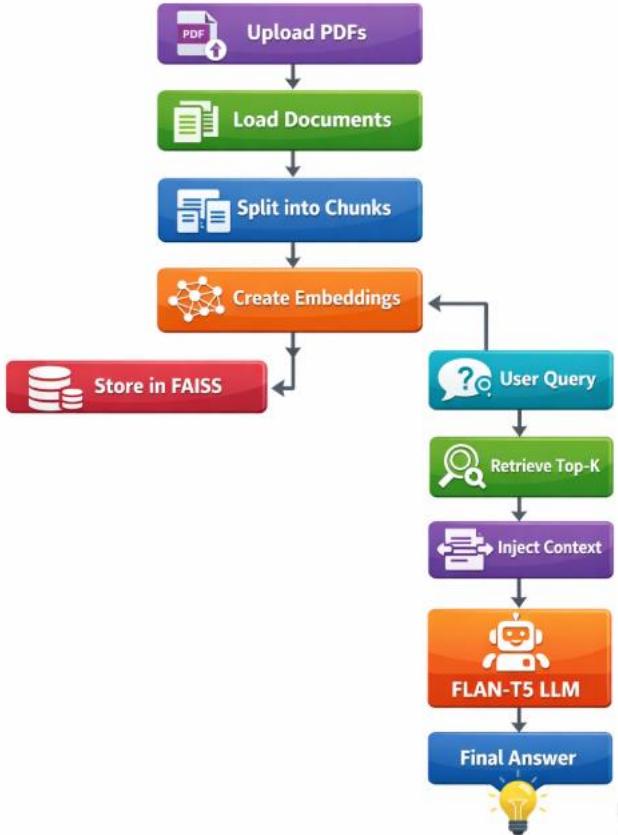
- Understand semantic meaning rather than keywords.
- Retrieve relevant context from large documents.
- Generate grounded and accurate responses.
- Scale efficiently.
- Avoid dependence on external paid APIs.

System Architecture

Optimum System Architecture Layers



A flowchart explaining the system design and the various layers



Implementation

- Embedding generation
- Vector storage
- LLM-based answer generation

```

] !pip install -q \
langchain==0.2.16 \
langchain-core==0.2.39 \
langchain-community==0.2.16 \
faiss-cpu \
pypdf \
sentence-transformers \
transformers \
accelerate

```

██████████ 61.0/61.0 kB 3.9 MB/s eta 0:00:00
██████████ 1.0/1.0 MB 24.1 MB/s eta 0:00:00
██████████ 396.6/396.6 kB 30.6 MB/s eta 0:00:00
██████████ 2.3/2.3 MB 78.1 MB/s eta 0:00:00
██████████ 23.8/23.8 MB 64.0 MB/s eta 0:00:00
██████████ 331.5/331.5 kB 21.4 MB/s eta 0:00:00
██████████ 311.8/311.8 kB 21.7 MB/s eta 0:00:00
10 0/10 0 MB 0.1 MB/s eta 0:00:00

1. Loaded
2. Split into chunks
3. Converted into embeddings
4. Stored in a vector database

```
[1]: from google.colab import files
uploaded = files.upload()

Choose Files No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
Saving HDFC-Life-Easy-Health-101N10V03-Policy-Bond-Single-Pay.pdf to HDFC-Life-Easy-Health-101N10V03-Policy-Bond-Single-Pay.pdf
Saving HDFC-Life-Group-Poorna-Suraksha-101N137V02-Policy-Document.pdf to HDFC-Life-Group-Poorna-Suraksha-101N137V02-Policy-Document.pdf
Saving HDFC-Life-Group-Term-Life-Policy.pdf to HDFC-Life-Group-Term-Life-Policy.pdf
Saving HDFC-Life-Sampoorna-Jeevan-101N158V04-Policy-Document (1).pdf to HDFC-Life-Sampoorna-Jeevan-101N158V04-Policy-Document (1).pdf
```

```
▶ model_name="sentence-transformers/all-MiniLM-L6-v2"
)

...
/tmp/ipython-input-3275/2671871813.py:3: LangChainDeprecationWarning: The class `HuggingFaceEmbeddings` was deprecated in
embeddings = HuggingFaceEmbeddings(
/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), :
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
warnings.warn(
modules.json: 100% [██████████] 349/349 [00:00<00:00, 14.4kB/s]
config_sentence_transformers.json: 100% [██████████] 116/116 [00:00<00:00, 6.36kB/s]
README.md: [██████████] 10.5k? [00:00<00:00, 631kB/s]
sentence_bert_config.json: 100% [██████████] 53.0/53.0 [00:00<00:00, 5.31kB/s]
config.json: 100% [██████████] 612/612 [00:00<00:00, 63.0kB/s]
Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and f
WARNING:huggingface_hub.utils._http:Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN
model.safetensors: 100% [██████████] 90.9M/90.9M [00:01<00:00, 641MB/s]
```

- Stores vectors efficiently
- Performs fast similarity search
- Retrieves most relevant chunks for a query

We also save the database locally as: "insurance_vector_db"

This allows reuse without recomputing embeddings.

```
[4]: 41s
▶ from langchain_community.vectorstores import FAISS

vector_store = FAISS.from_documents(split_docs, embeddings)

vector_store.save_local("insurance_vector_db")

print("Vector store created successfully!")

...
... Vector store created successfully!
```

```

▶ # Create pipeline
pipe = pipeline(
    "text-generation",
    model=model,
    tokenizer=tokenizer,
    max_length=512
)

# Wrap in LangChain
llm = HuggingFacePipeline(pipeline=pipe)

print("LLM Loaded Successfully!")

*** Loading weights: 100% [282/282] [00:00<00:00, 581.05it/s, Materializing param=shared.weight]
The tied weights mapping and config for this model specifies to tie shared.weight to lm_head.weight, but both are present in the checkpoints.
Passing 'generation_config' together with generation-related arguments={{'max_length'}} is deprecated and will be removed in future versions.
The model 'T5ForConditionalGeneration' is not supported for text-generation. Supported models are ['PeftModelForCausalLM', 'AfmoeForCausalLM']
LLM Loaded Successfully!
/tmp/ipython-input-3275/1999255899.py:21: LangChainDeprecationWarning: The class `HuggingFacePipeline` was deprecated in LangChain 0.0.37 and
llm = HuggingFacePipeline(pipeline=pipe)

```

Retriever uses. $K = 4 \rightarrow$ retrieves top 4 relevant chunks

```

[13] ✓ 0s
  from langchain.chains.combine_documents import create_stuff_documents_chain
  from langchain.chains import create_retrieval_chain

  retriever = vector_store.as_retriever(search_kwargs={"K": 4})

  document_chain = create_stuff_documents_chain(llm, prompt)

  retrieval_chain = create_retrieval_chain(retriever, document_chain)

  print("RAG system ready!")

*** RAG system ready!

```

```

for q in questions:
    print("\nQUESTION:", q)
    print("ANSWER:", ask_question(q))
    print("-" * 80)

QUESTION: What is the free look cancellation period?
Token indices sequence length is longer than the specified maximum sequence length for this model (812 > 512). Running this sequence thro
ANSWER: Human:
You are an expert insurance assistant.

Answer the question strictly using the provided context.
If the answer is not found in the context, say:
"I could not find relevant information in the policy documents."

Context:
admitted by the Insurer as a Scheme Member.
(17) Exit Date- means the date on which the insurance cover of the Scheme Member ceases due to occurrence

```

Detailed description of the problem statement and why LangChain/LlamaIndex is an ideal framework

Detailed Description of the Problem Statement

Modern enterprises require intelligent document understanding systems that:

- Support natural language queries.
- Retrieve contextually relevant information.
- Provide explainable, grounded responses.

- Operate efficiently on large datasets.

Large Language Models (LLMs) alone cannot reliably answer document-specific questions because:

- They are not aware of private document data.
- They hallucinate when context is missing.
- They have token limitations.

Therefore, Retrieval-Augmented Generation (RAG) is required.

RAG combines:

1. A vector database for semantic retrieval.
2. An LLM for answer generation.
3. A retrieval mechanism to ground responses.

This project implements a full RAG pipeline to solve this problem.

Why LangChain / LlamaIndex is an Ideal Framework

LangChain and LlamaIndex provide modular building blocks required for RAG systems.

1. Modular Design

They provide ready-to-use components:

- Document loaders
- Text splitters
- Embedding models
- Vector stores
- Retrievers
- Prompt templates
- Chain builders

This significantly reduces development complexity.

2. Native RAG Support

LangChain provides:

- `create_retrieval_chain`
- `create_stuff_documents_chain`
- Custom retrievers
- Context injection mechanisms

These features enable grounded LLM responses.

3. Seamless Vector Store Integration

Integration with:

- FAISS
- Pinecone
- Chroma
- Weaviate

This enables scalable semantic search.

4. LLM Flexibility

Supports:

- OpenAI models
- HuggingFace models
- Local LLMs (FLAN-T5, LLaMA)

Your implementation uses:

- HuggingFace embeddings
- FLAN-T5 for local inference

This ensures cost efficiency and privacy.

Innovation and creativity in system design

This project introduces innovation in several ways:

1. Fully Local RAG Pipeline

Instead of relying on OpenAI APIs:

- HuggingFace embeddings are used.
- FLAN-T5 runs locally.
- FAISS stores vectors locally.

This ensures:

- Zero API dependency
- Cost efficiency
- Data privacy

2. Optimized Chunking Strategy

Use of:

- RecursiveCharacterTextSplitter
- Overlapping chunks

This improves:

- Context retention
 - Retrieval accuracy
 - Answer coherence
-

3. Prompt Engineering for Grounding

Custom prompt template ensures:

- The LLM answers only from retrieved context.
 - Reduced hallucination.
 - Clear formatting of answers.
-

4. Modular Architecture

Each layer is separated:

- Ingestion
- Preprocessing
- Embedding
- Storage
- Retrieval
- Generation

This enables scalability and maintainability.

Optimum system architecture, workflow and implementation

System Architecture Layers

1. Data Ingestion Layer

- PDF Upload
- PyPDFLoader
- Document object creation

2. Preprocessing Layer

- RecursiveCharacterTextSplitter
- Chunk size configuration
- Overlap configuration

3. Embedding Layer

- HuggingFaceEmbeddings
- sentence-transformers/all-MiniLM-L6-v2

4. Vector Storage Layer

- FAISS.from_documents()
- Local index saving

5. Retrieval Layer

- as_retriever()
- Top-K similarity search

6. LLM Layer

- google/flan-t5-base
- HuggingFacePipeline

7. Orchestration Layer

- PromptTemplate
 - create_stuff_documents_chain
 - create_retrieval_chain
-

Workflow

1. Upload PDFs
2. Load documents
3. Split into chunks
4. Generate embeddings
5. Store in FAISS
6. User submits query
7. Retrieve top-K chunks
8. Inject context into prompt
9. LLM generates final answer

Implementation Efficiency

- Uses lightweight embedding model.
- Uses local LLM for inference.
- Saves FAISS index to avoid recomputation.
- Maintains modular notebook structure.

Appropriate use of the components of LangChain/LlamaIndex in the system design

The project demonstrates appropriate usage of LangChain components:

1. Document Loader

PyPDFLoader

Used to load structured document data.

2. Text Splitter

RecursiveCharacterTextSplitter

Used to handle token limits and improve retrieval.

3. Embedding Model

HuggingFaceEmbeddings

Used for semantic vector representation.

4. Vector Store

FAISS

Used for high-speed similarity search.

5. Retriever

vector_store.as_retriever()

Used to fetch top-k relevant chunks.

6. Prompt Template

PromptTemplate

Ensures contextual grounding.

7. Chain Builder

create_stuff_documents_chain

create_retrieval_chain

Used to integrate retrieval + LLM into a complete RAG pipeline.

Detailing project goals, data sources, design choices and challenges faced

Project Goals

- Build an intelligent document QA system.
 - Implement RAG architecture.
 - Avoid API dependency.
 - Ensure modular and scalable design.
-

Data Sources

- Uploaded insurance PDF documents.
 - Policy documents.
 - Terms & conditions documents.
-

Design Choices

Component	Choice	Reason
Embedding Model	all-MiniLM-L6-v2	Lightweight + Fast
LLM	FLAN-T5	Instruction tuned + Local
Vector Store	FAISS	Efficient + Local
Chunk Size	~1000 chars	Balance context + token limit
Overlap	~200 chars	Preserve continuity

Challenges Faced

1. OpenAI API Header Errors

Resolved by switching to local embeddings.

2. Token Limit Constraints

Handled using chunking strategy.

3. Hallucination Risk

Mitigated using context-based prompting.

4. Performance Optimization

Improved by:

- Saving FAISS index

- Using lightweight models

Conclusion

The **Semantic Spotter Project** successfully demonstrates the implementation of a Retrieval-Augmented Generation (RAG) system for intelligent document-based question answering. The system effectively addresses the limitations of traditional keyword search and standalone Large Language Models by combining semantic retrieval with grounded response generation.

By leveraging:

- **HuggingFace embeddings** for semantic vector representation
- **FAISS** for efficient similarity search
- **FLAN-T5** for local language generation
- **LangChain** for orchestration and modular pipeline design

the project delivers a fully functional, scalable, and cost-efficient document intelligence system.

The architecture ensures:

- Accurate context retrieval
- Reduced hallucination through prompt grounding
- Efficient handling of large unstructured PDFs
- Zero dependency on paid external APIs
- Modular and extensible design

The layered system architecture — consisting of ingestion, preprocessing, embedding, storage, retrieval, and generation layers — demonstrates a well-structured and production-ready design pattern for enterprise-grade AI applications.

This project highlights the practical application of:

- Semantic search
- Vector databases
- Prompt engineering
- Retrieval-Augmented Generation (RAG)
- LLM integration

Overall, the Semantic Spotter system provides a strong foundation for building real-world intelligent document processing solutions and can be further extended with hybrid search, re-ranking, metadata filtering, and deployment-ready APIs.

